

Task quality comparison

Print illustration · Longer task and evidence records are excerpted.

User evaluation demand

I want to evaluate whether a model can satisfy multiple instructions and constraints that are simultaneous, compositional, and out-of-distribution when serving as an AI assistant.

A

Candidate task

Free response

▼ Task

Output just with the final answer to the question.
Question:Which rare-earth subgroup has ore bodies that are more rare, smaller, and less concentrated?
Answer:

▼ Recorded reference answer

HREE

▼ Recorded model response

Heavy rare-earth elements (HREE)

B

Candidate task

Multi-turn

▼ Task

Write an incident timeline for last night's payment disruption. Reply with the timeline and nothing else.

Format: exactly five numbered entries, one per line, each exactly 'N.
[Excerpt; continued in the full record.]

► Scoring rubric (excerpt)

► Dialogue policy and all turns

▼ Recorded model response

User (turn 3; excerpt):
Add a sixth entry at 14:04 as the last numbered item, with its own text. Where that clashes with anything I said earlier, the new value wins for the new
[Excerpt; continued in the full record.]

Assistant (turn 3; excerpt):
[Entries 1-4 omitted here.]
5. 13:55 - Engineering restored normal payment flow.
6. 14:04 - Fixes verified across transaction channels.
Even-minute entries: 5

Which task is better?

Choose the task that better evaluates this demand. Use the recorded response as supporting evidence.

☐ Prefer A

☐ Prefer B

Reason (optional)

Submit and continue